

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
%matplotlib inline
```

```
In [2]: from sklearn.datasets import load_boston

# Load the Boston housing dataset
boston_dataset = load_boston()

# Print the keys of the dataset
print(boston_dataset.keys())
```

-----  
**ImportError**

Traceback (most recent call last)

Cell In[2], line 1

```
----> 1 from sklearn.datasets import load_boston
      3 # Load the Boston housing dataset
      4 boston_dataset = load_boston()
```

File ~\anaconda3\lib\site-packages\sklearn\datasets\\_\_init\_\_.py:156, in \_\_getattr\_\_(name)

```
105 if name == "load_boston":
106     msg = textwrap.dedent(
107         """
108         `load_boston` has been removed from scikit-learn since version 1.2.
109     (...)
154         """
155     )
--> 156     raise ImportError(msg)
157 try:
158     return globals()[name]
```

**ImportError:**

`load\_boston` has been removed from scikit-learn since version 1.2.

The Boston housing prices dataset has an ethical problem: as investigated in [1], the authors of this dataset engineered a non-invertible variable "B" assuming that racial self-segregation had a positive impact on house prices [2]. Furthermore the goal of the research that led to the creation of this dataset was to study the impact of air quality but it did not give adequate demonstration of the validity of this assumption.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

[1] M Carlisle.  
"Racist data destruction?"  
<<https://medium.com/@docintangible/racist-data-destruction-113e3eff54a8>>

[2] Harrison Jr, David, and Daniel L. Rubinfeld.

"Hedonic housing prices and the demand for clean air."  
Journal of environmental economics and management 5.1 (1978): 81-102.  
<[https://www.researchgate.net/publication/4974606\\_Hedonic\\_housing\\_prices\\_and\\_the\\_demand\\_for\\_clean\\_air](https://www.researchgate.net/publication/4974606_Hedonic_housing_prices_and_the_demand_for_clean_air)>

In [ ]:

In [3]: dataset = pd.read\_csv("C:\Users\omkar\Downloads\boston.csv")

Cell In[3], line 1

dataset = pd.read\_csv("C:\Users\omkar\Downloads\boston.csv")

**SyntaxError:** (unicode error) 'unicodeescape' codec can't decode bytes in position 2-3: truncated \UXXXXXXX escape

In [4]: dataset = pd.read\_csv("C:/Users/omkar/Downloads/boston.csv")

In [5]: dataset.keys()

Out[5]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MDEV'], dtype='object')

In [6]: dataset.head()

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MDEV
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

In [7]: dataset.describe()

Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	B	LSTAT	MDEV
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.593761	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549801	15.23	392.82
std	8.596783	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707119	4.037563	34.562374
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	4.98	21.6
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	15.23	392.82
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	15.23	392.82
75%	3.647423	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	15.23	392.82
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	15.23	392.82

In [8]: dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM        506 non-null    float64
1    ZN          506 non-null    float64
2    INDUS       506 non-null    float64
3    CHAS        506 non-null    float64
4    NOX         506 non-null    float64
5    RM          506 non-null    float64
6    AGE         506 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    float64
9    TAX         506 non-null    float64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       506 non-null    float64
13   MDEV        506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

```
In [9]: dataset.isnull().sum()
```

```
Out[9]: CRIM      0
        ZN       0
        INDUS   0
        CHAS    0
        NOX     0
        RM      0
        AGE     0
        DIS     0
        RAD     0
        TAX     0
        PTRATIO 0
        B       0
        LSTAT   0
        MDEV    0
dtype: int64
```

```
In [10]: dataset = dataset.fillna(dataset.mean())
```


```
In [11]: dataset.isnull().sum()
```

```
Out[11]: CRIM      0
        ZN       0
        INDUS   0
        CHAS    0
        NOX     0
        RM      0
        AGE     0
        DIS     0
        RAD     0
        TAX     0
        PTRATIO 0
        B       0
        LSTAT   0
        MDEV    0
dtype: int64
```

```
In [12]: #Plot the distribution of 'MDEV' = median value of owner-occupied homes in thousands of
```

```
In [13]: sns.set(rc={'figure.figsize':(11.7,8.27)})
        sns.displot(dataset['MDEV'],bins=30);
```

```
plt.show()
```



```
In [14]: #correlation matrix
```

```
In [15]: correlation_matrix = dataset.corr().round(2)
sns.heatmap(data=correlation_matrix, annot=True)
```

```
Out[15]: <Axes: >
```



```
In [16]: plt.figure(figsize=(20, 5))

# Define the features and the target variable
features = ['LSTAT', 'RM']
target = dataset['MDEV']



# Loop through each feature
for i, col in enumerate(features):
    # Create subplots
    plt.subplot(1, len(features), i + 1)

    # Define x and y values for the scatter plot
    x = dataset[col]
    y = target

    # Plot the scatter plot
    plt.scatter(x, y, marker='o')

    # Set title, xlabel, and ylabel for the subplot
    plt.title(col)
    plt.xlabel(col)
    plt.ylabel('MDEV')

plt.show()
```



```
In [17]: import numpy as np
x = pd.DataFrame(np.c_[dataset['LSTAT'], dataset['RM']], columns=['LSTAT', 'RM'])
y = dataset['MDEV']
```

```
In [18]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=5)

print("Training set shapes:")
print("x_train:", x_train.shape)
print("y_train:", y_train.shape)
print("Testing set shapes:")
print("x_test:", x_test.shape)
print("y_test:", y_test.shape)

Training set shapes:
x_train: (404, 2)
y_train: (404,)
Testing set shapes:
x_test: (102, 2)
y_test: (102,)
```

```
In [19]: from sklearn.linear_model import LinearRegression
model=LinearRegression()

model.fit(x_train,y_train)
```

```
Out[19]: ▼ LinearRegression
LinearRegression()
```

```
In [20]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
import numpy as np

y_pred = model.predict(x_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))

r2 = r2_score(y_test, y_pred)

print("Model performance for testing set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))

Model performance for testing set
-----
RMSE is 5.137400784702911
.6628996975186953
```

In [21]: *# Predicting selling price*

```
sample_data = [[6.89, 9.939]]
```

```
price = model.predict(sample_data)
```

```
print("Predicted selling price for house: ${:,.2f}".format(price[0]))
```

Predicted selling price for house: \$43.41

C:\Users\omkar\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(

In [22]: `sample_data_df = pd.DataFrame(sample_data, columns=['LSTAT', 'RM'])`

```
price = model.predict(sample_data_df)
```

```
print("Predicted selling price for house: ${:,.2f}".format(price[0]))
```

Predicted selling price for house: \$43.41

In [ ]: